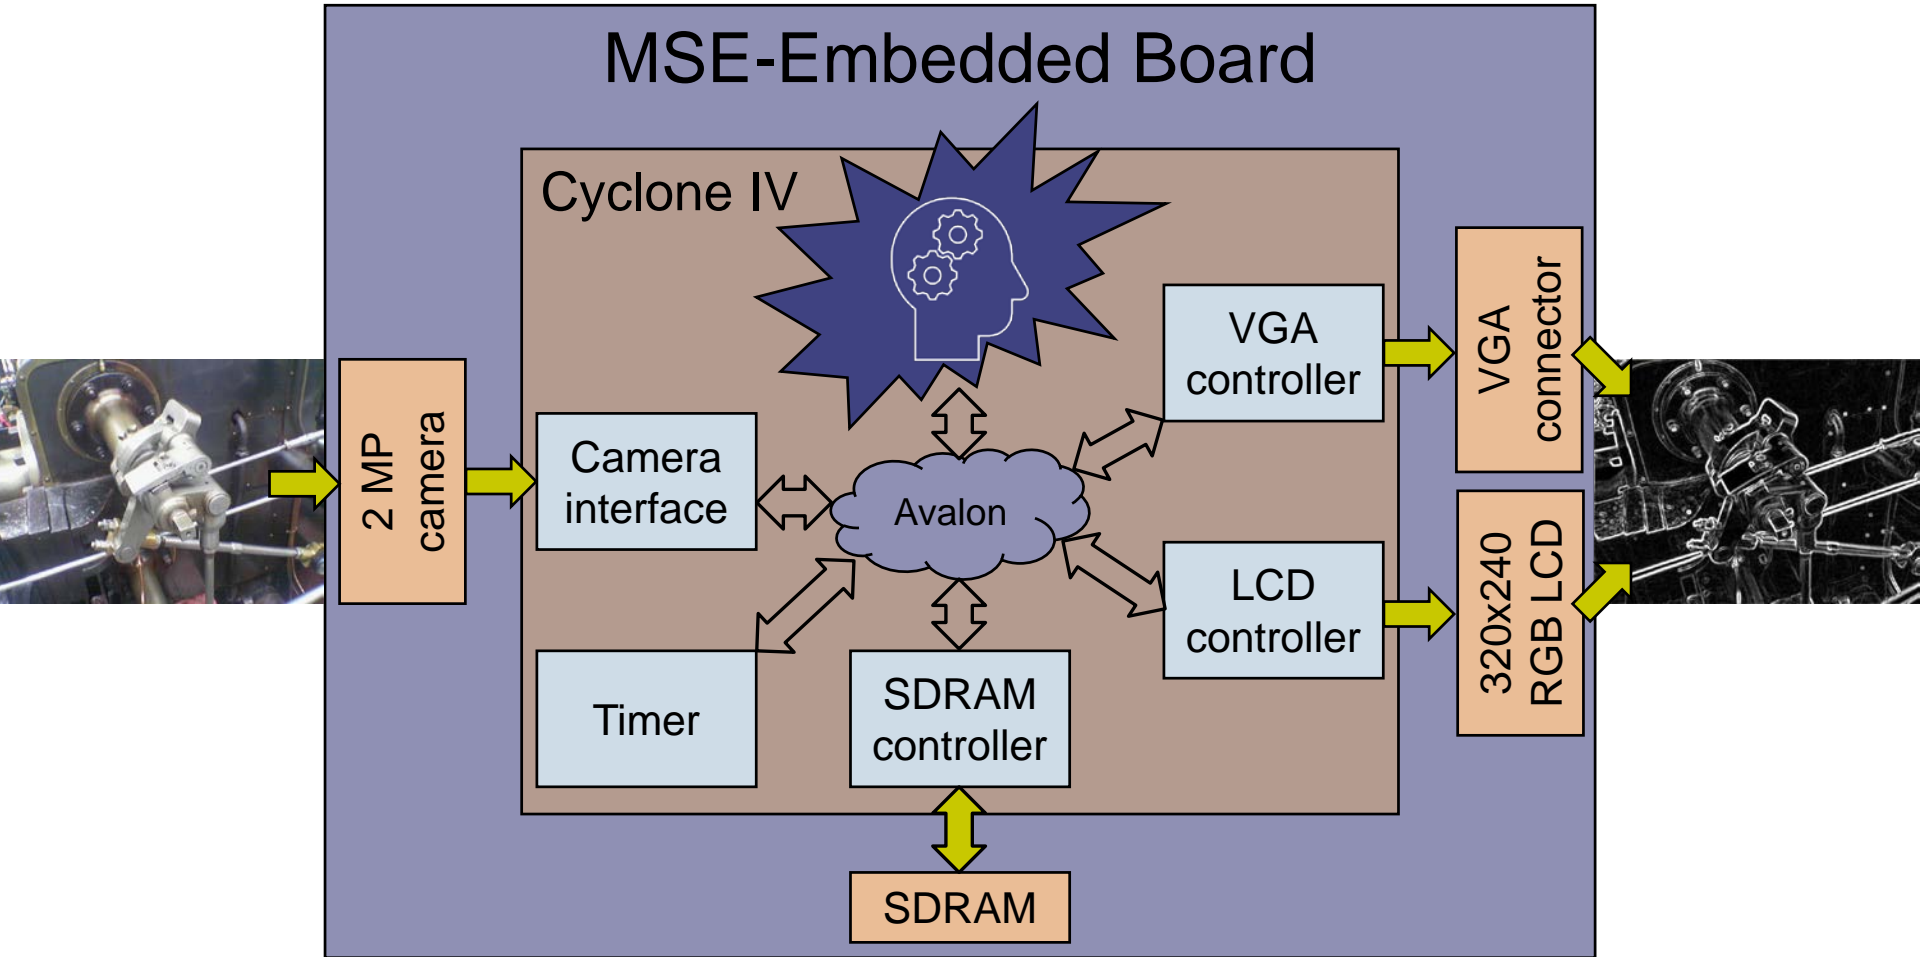


Design of Embedded Hardware and Firmware Software Optimization 2

Andrea Guerrieri
HES-SO//Genève
andrea.guerrieri@hesge.ch

Problem specifications



«Out-of-the-box» optimization

Function	-O0	-O1	-O2	-O3
conv_grayscale (Time)	127336549 (2.547 s)	44161114 (0.883 s)	41296763 (0.826 s)	41194172 (0.824 s) 3.09x
sobel_x (Time)	729624798 (14.592 s)	129966322 (2.599 s)	109616923 (2.192 s)	29727206 (0.595 s) 24.54x
sobel_y (Time)	729677739 (14.594 s)	129964177 (2.599 s)	109607580 (2.192 s)	27249940 (0.545 s) 26.78x
sobel_threshold (Time)	102849150 (2.057 s)	33664258 (0.673 s)	32015554 (0.64 s)	32126201 (0.643 s) 3.20x
TOTAL (Time)	1689488236 (33.79 s)	337755871 (6.754 s)	292536820 (5.85 s)	130297519 (2.607 s) 12.97x

Loop unrolling + in-lining


Function	-O0	-O0 (no loops)	-O0 (no loops, inline)
conv_grayscale (Time)	127336549 (2.547 s)	127071260 (2.541 s)	125935754 (2.519 s)
sobel_x (Time)	729624798 (14.592 s)	425314676 (8.506 s)	325845850 (6.517 s)
sobel_y (Time)	729677739 (14.594 s)	425378323 (8.508 s)	320545780 (6.411 s)
sobel_threshold (Time)	102849150 (2.057 s)	105635977 (2.113 s)	104956605 (2.099 s)
TOTAL (Time)	1689488236 (33.79 s)	1083400236 (21.67 s)	877283989 (17.55 s)

1.31x
(2.24x)

1.56x

1.93x

Loop unrolling + in-lining



Software Optimizations

So far only code-level modifications to remove the overhead...

Software Optimizations

So far only code-level modifications to remove the overhead...

How can we improve more?


RGB to grayscale

➤ First approximation:

$$\text{Grayscale} = (R+G+B)/3$$

RGB to grayscale

- First approximation:
Grayscale = $(R+G+B)/3$
- Rather black output



RGB to grayscale

- First approximation:

$$\text{Grayscale} = (R+G+B)/3$$

- Rather black output



- Luminance (ITU-R Rec.601), approximated:

$$\text{Grayscale} = 0.3*R + 0.59*G + 0.11*B$$


RGB to grayscale

$$\text{Grayscale} = 0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B$$

```
void conv_grayscale(void *picture,
                    int width,
                    int height) {
    int x,y,gray;
    unsigned short *pixels = (unsigned short *)picture , rgb;
    unsigned int temp;
    grayscale_width = width;
    grayscale_height = height;
    if (grayscale_array != NULL)
        free(grayscale_array);
    grayscale_array = (unsigned char *) malloc(width*height);
    for (y = 0 ; y < height ; y++) {
        for (x = 0 ; x < width ; x++) {
            rgb = pixels[y*width+x];
            temp = (rgb>>11)&0x1F; // red value
            gray = (temp*30)/100;
            temp = (rgb>>5)&0x3F; // green value
            gray += (temp*59)/100;
            temp = rgb&0x1F; // blue value
            gray += (temp*11)/100;
            IOWR_8DIRECT(grayscale_array,y*width+x,gray);
        }
    }
}
```


Data Flow Graph (DFG)

$$\text{Grayscale} = 0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B$$



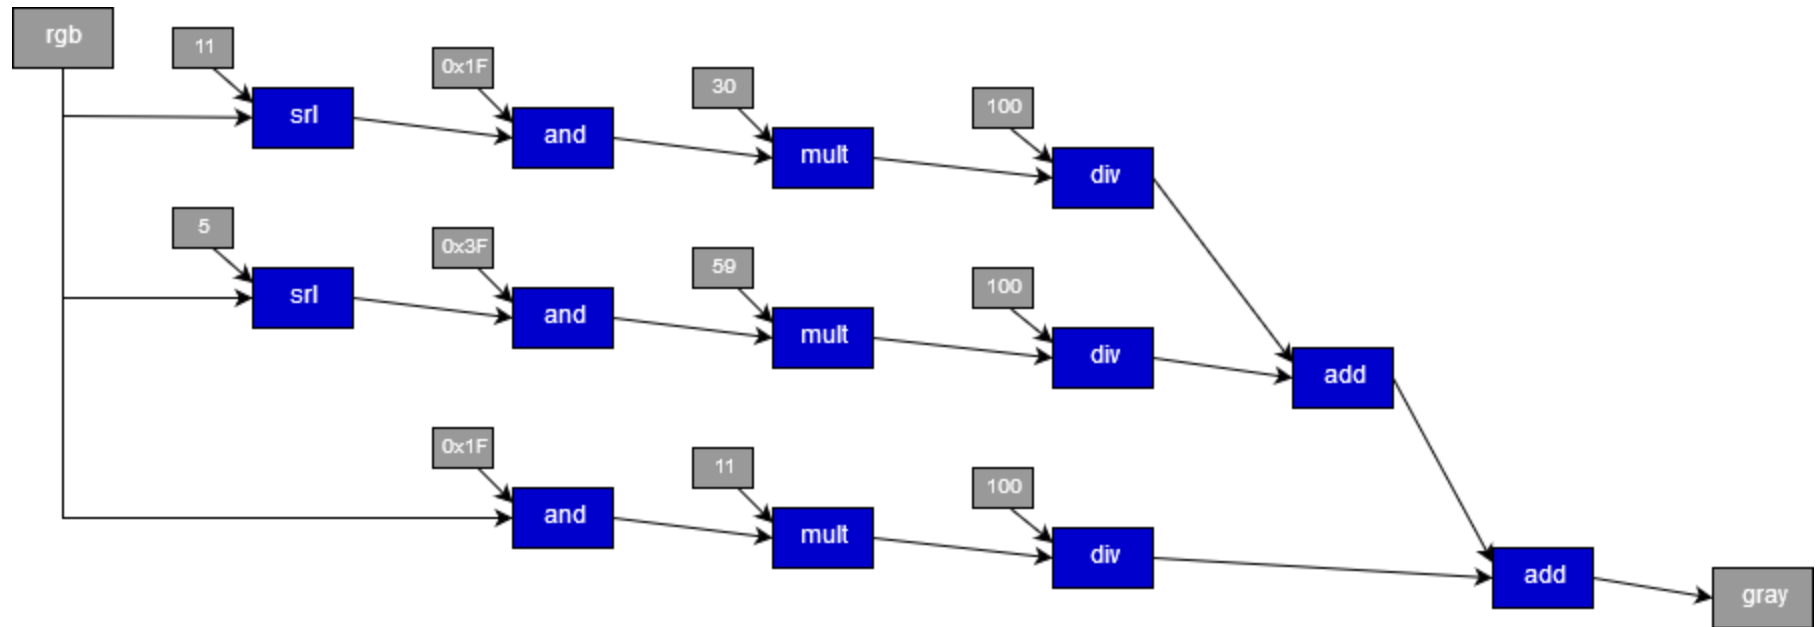
Data Flow Graph (DFG)

$$\text{Grayscale} = 0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B$$



Data Flow Graph (DFG)

$$\text{Grayscale} = 0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B$$



RGB to grayscale

$$\text{Grayscale} = 0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B$$

```
void conv_grayscale(void *picture,
                    int width,
                    int height) {
    int x,y,gray;
    unsigned short *pixels = (unsigned short *)picture , rgb;
    unsigned int temp;
    grayscale_width = width;
    grayscale_height = height;
    if (grayscale_array != NULL)
        free(grayscale_array);
    grayscale_array = (unsigned char *) malloc(width*height);
    for (y = 0 ; y < height ; y++) {
        for (x = 0 ; x < width ; x++) {
            rgb = pixels[y*width+x];
            temp = (rgb>>11)&0x1F; // red value
            gray = (temp*30)/100;
            temp = (rgb>>5)&0x3F; // green value
            gray += (temp*59)/100;
            temp = rgb&0x1F; // blue value
            gray += (temp*11)/100;
            IOWR_8DIRECT(grayscale_array,y*width+x,gray);
        }
    }
}
```

~945 cycles/pixel

RGB to grayscale

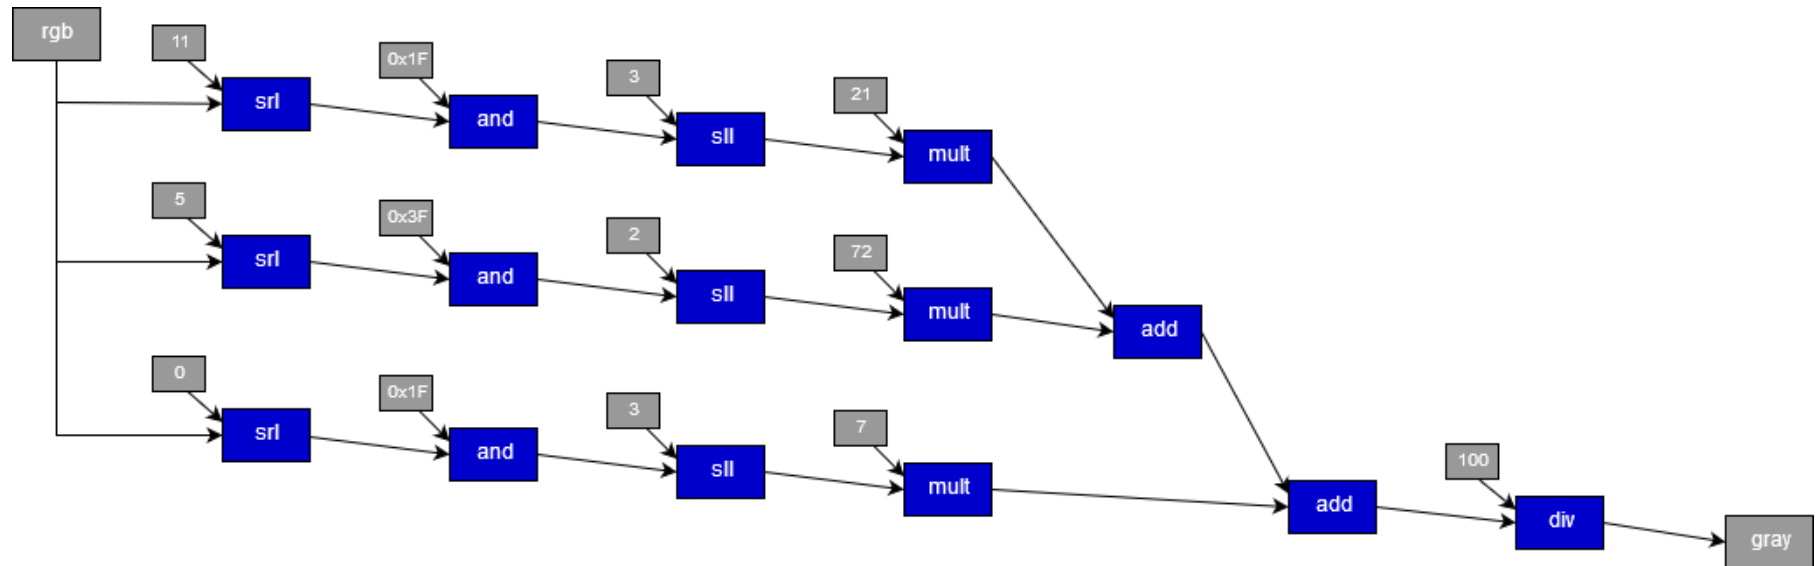
$$\text{Grayscale} \approx 0.33 \cdot R + 0.56 \cdot G + 0.11 \cdot B$$

```
void conv_grayscale(void *picture,
                    int width,
                    int height) {
    int x,y,gray;
    unsigned short *pixels = (unsigned short *)picture , rgb;
    grayscale_width = width;
    grayscale_height = height;
    if (grayscale_array != NULL)
        free(grayscale_array);
    grayscale_array = (unsigned char *) malloc(width*height);
    for (y = 0 ; y < height ; y++) {
        for (x = 0 ; x < width ; x++) {
            rgb = pixels[y*width+x];
            gray = (((rgb>>11)&0x1F)<<3)*21; // red part
            gray += (((rgb>>5)&0x3F)<<2)*72; // green part
            gray += (((rgb>>0)&0x1F)<<3)*7; // blue part
            gray /= 100;
            IOWR_8DIRECT(grayscale_array,y*width+x,gray);
        }
    }
}
```

~800 cycles/pixel

Data Flow Graph (DFG)

$$\text{Grayscale} \approx 0.33 * R + 0.56 * G + 0.11 * B$$



Algorithm optimization

Objdump:

```

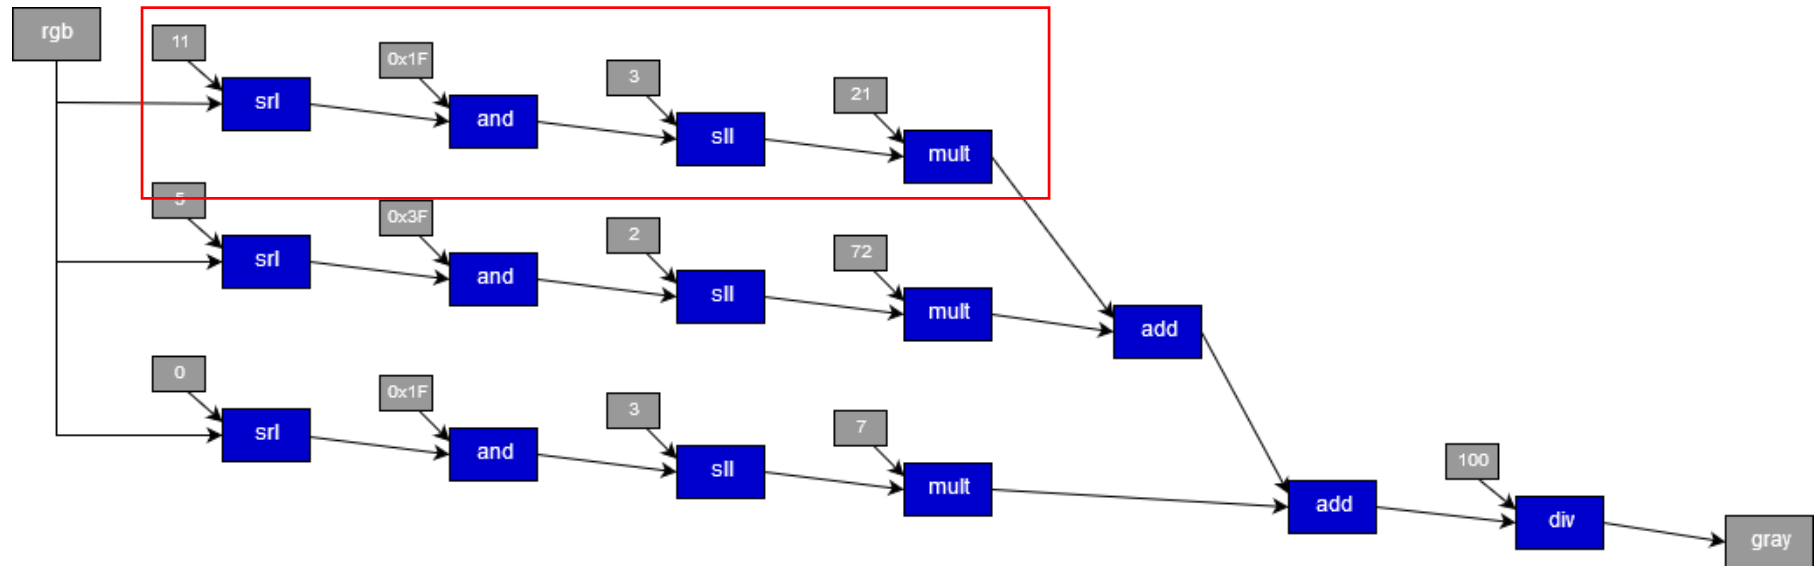
b48: 1080000b ldhu    r2,0(r2)
b4c: e0bffb0d sth     r2,-20(fp)
        gray = (((rgb>>11)&0x1F)<<3)*21; // red part
b50: e0bffb0b ldhu    r2,-20(fp)
b54: 1004d2fa srli    r2,r2,11
b58: 10bfffcc andi    r2,r2,65535
b5c: 100490fa slli    r2,r2,3
b60: 10803fcc andi    r2,r2,255
b64: 10800564 muli    r2,r2,21
b68: e0bffc15 stw     r2,-16(fp)
        gray += (((rgb>>5)&0x3F)<<2)*72; // green part
b6c: e0bffb0b ldhu    r2,-20(fp)
b70: 1004d17a srli    r2,r2,5
b74: 10bfffcc andi    r2,r2,65535
b78: 1085883a add     r2,r2,r2
b7c: 1085883a add     r2,r2,r2
b80: 10803fcc andi    r2,r2,255
b84: 10801224 muli    r2,r2,72
b88: e0fffc17 ldw     r3,-16(fp)
b8c: 1885883a add     r2,r3,r2
b90: e0bffc15 stw     r2,-16(fp)
        gray += (((rgb>>0)&0x1F)<<3)*7; // blue part
b94: e0bffb0b ldhu    r2,-20(fp)
b98: 100490fa slli    r2,r2,3
b9c: 10803fcc andi    r2,r2,255
ba0: 108001e4 muli    r2,r2,7
ba4: e0fffc17 ldw     r3,-16(fp)
ba8: 1885883a add     r2,r3,r2
bac: e0bffc15 stw     r2,-16(fp)
        gray /= 100;
bb0: e0bffc17 ldw     r2,-16(fp)
bb4: 01401904 movi    r5,100
bb8: 1009883a mov     r4,r2
bbc: 00022c80 call   22c8 <__divsi3>
bc0: e0bffc15 stw     r2,-16(fp)
        IOWR_8DIRECT(gray, grayscale_array, y*width+x, gray);
bc4: d0a6cd17 ldw     r2,-25804(gp)
bc8: e13ff917 ldw     r4,-28(fp)

```

Data Flow Graph (DFG)

$$\text{Grayscale} \approx 0.33 * R + 0.56 * G + 0.11 * B$$

Operations:
7



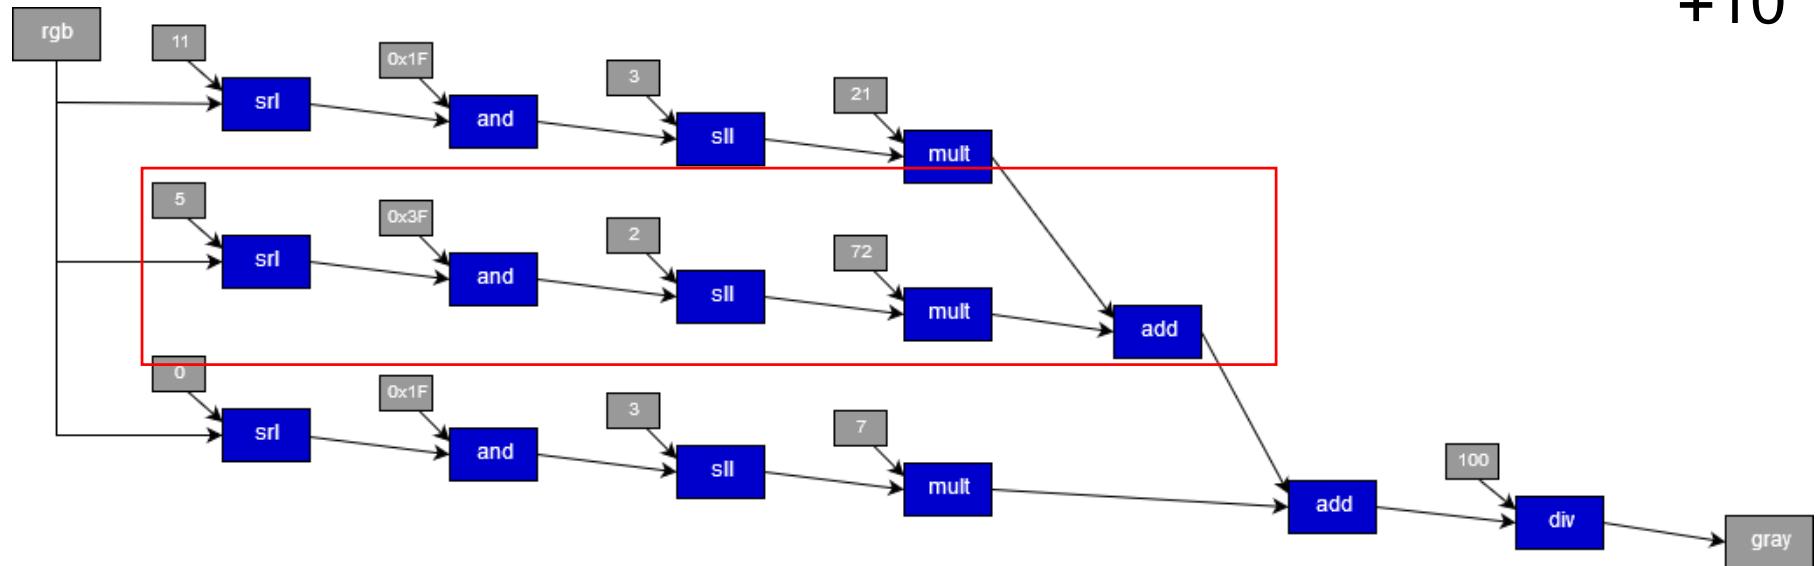
Data Flow Graph (DFG)

$$\text{Grayscale} \approx 0.33 * R + 0.56 * G + 0.11 * B$$

Operations:

7

+10



Data Flow Graph (DFG)

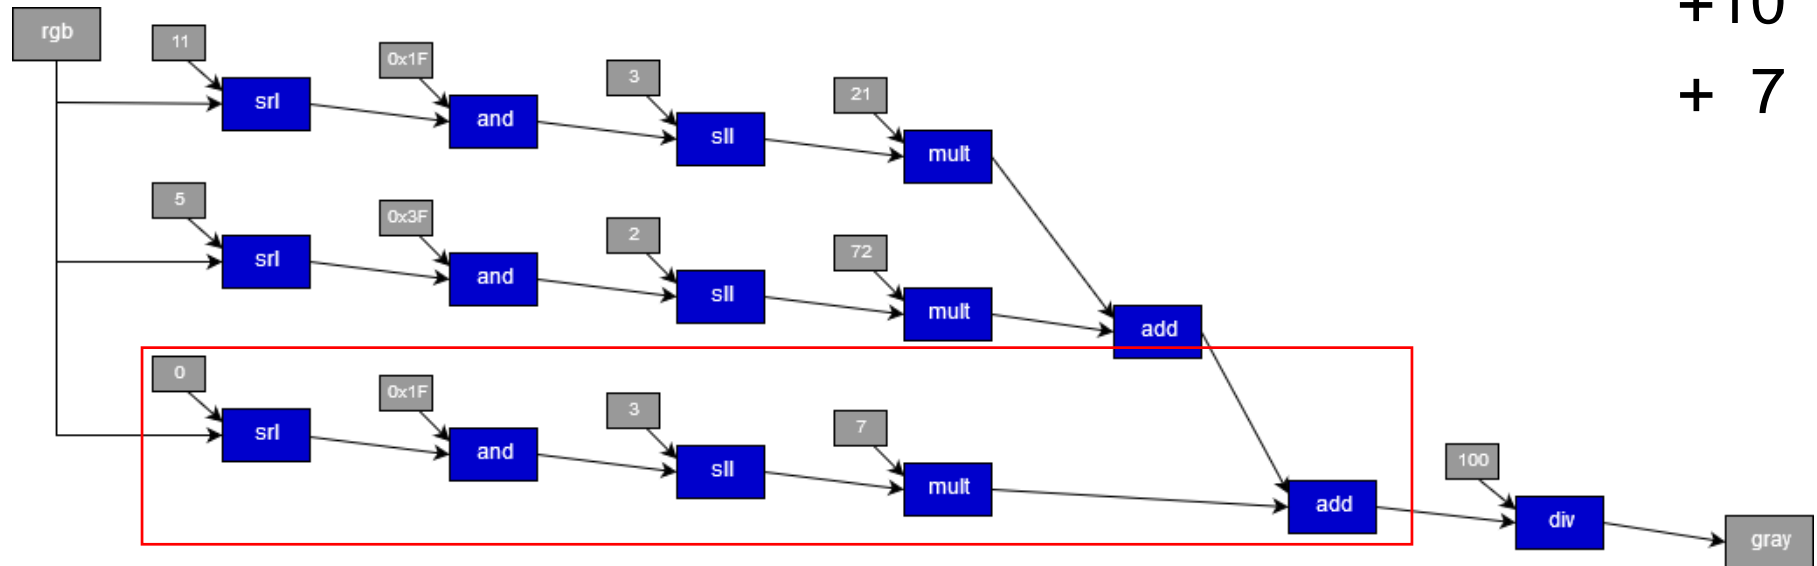
$$\text{Grayscale} \approx 0.33 \cdot R + 0.56 \cdot G + 0.11 \cdot B$$

Operations:

7

+10

+ 7



Data Flow Graph (DFG)

$$\text{Grayscale} \approx 0.33 \cdot R + 0.56 \cdot G + 0.11 \cdot B$$


Operations:

7

+10

+ 7

+37



Algorithm optimization

Division:
33 operations

```
000022c8 <__divsi3>:
22c8: 20001b16      blt    r4,zero,2338 <__divsi3+0x70>
22cc: 000f883a      mov    r7,zero
22d0: 28001616      blt    r5,zero,232c <__divsi3+0x64>
22d4: 200d883a      mov    r6,r4
22d8: 29001a2e      bgeu   r5,r4,2344 <__divsi3+0x7c>
22dc: 00800804      movi   r2,32
22e0: 00c00044      movi   r3,1
22e4: 00000106      br     22ec <__divsi3+0x24>
22e8: 10000d26      beq     r2,zero,2320 <__divsi3+0x58>
22ec: 294b883a      add    r5,r5,r5
22f0: 10bffc4       addi   r2,r2,-1
22f4: 18c7883a      add    r3,r3,r3
22f8: 293ffb36      bltu   r5,r4,22e8 <__divsi3+0x20>
22fc: 0005883a      mov    r2,zero
2300: 18000726      beq     r3,zero,2320 <__divsi3+0x58>
2304: 0005883a      mov    r2,zero
2308: 31400236      bltu   r6,r5,2314 <__divsi3+0x4c>
230c: 314dc83a      sub    r6,r6,r5
2310: 10c4b03a      or     r2,r2,r3
2314: 1806d07a      srli   r3,r3,1
2318: 280ad07a      srli   r5,r5,1
231c: 183ffa1e      bne     r3,zero,2308 <__divsi3+0x40>
2320: 38000126      beq     r7,zero,2328 <__divsi3+0x60>
2324: 0085c83a      sub    r2,zero,r2
2328: f800283a      ret
232c: 014bc83a      sub    r5,zero,r5
2330: 39c0005c      xori   r7,r7,1
2334: 003fe706      br     22d4 <__divsi3+0xc>
2338: 0109c83a      sub    r4,zero,r4
233c: 01c00044      movi   r7,1
2340: 003fe306      br     22d0 <__divsi3+0x8>
2344: 00c00044      movi   r3,1
2348: 003fee06      br     2304 <__divsi3+0x3c>
```

Nios II Reference Guide

Table 60. Hardware Multiply and Divide Details for the Nios II/f Core

ALU Option	Hardware Details	Cycles per Instruction	Result Latency Cycles	Supported Instructions
No hardware multiply or divide	Multiply and divide instructions generate an exception	–	–	None
Logic elements	ALU includes 32 x 4-bit multiplier	11	+2	mul, muli
<i>continued...</i>				
32-bit multiplier	ALU includes 32 x 32-bit multiplier	1	+2	mul, muli, mulxss, mulxsu, mulxuu
16-bit multiplier	ALU includes 3 16 x 16-bit multiplier	1	+2	mul, muli
16-bit multiplier	ALU includes 4 16 x 16-bit multiplier	2	+2	mul, muli, mulxss, mulxsu, mulxuu
Hardware divide	ALU includes SRT Radix-2 divide circuit	35	+2	div, divu

Data Flow Graph (DFG)

$$\text{Grayscale} \approx 0.33 \cdot R + 0.56 \cdot G + 0.11 \cdot B$$

Operations:

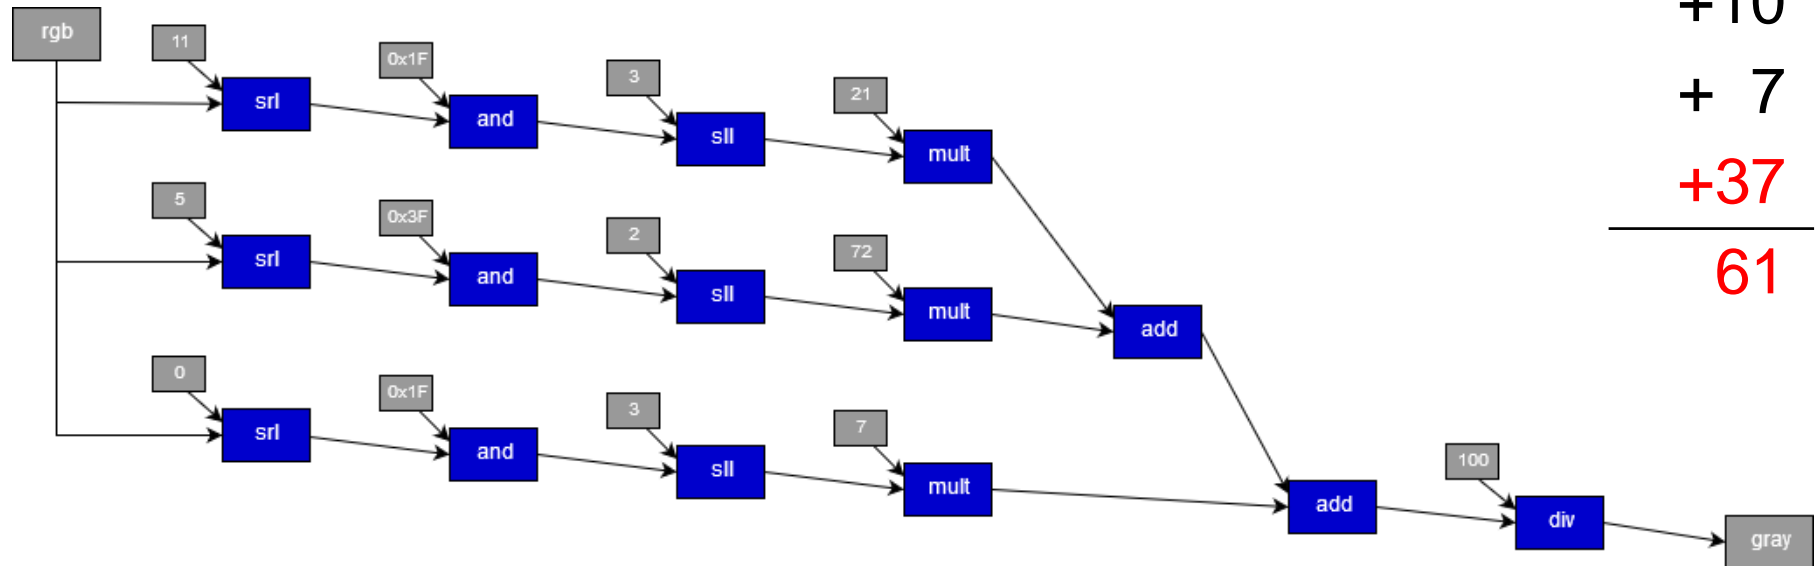
7

+10

+ 7

+37

61



Algorithm optimization

- Do we need a lot of precision in the division?

Algorithm optimization

- Do we need a lot of precision in the division?
- Not really, we could substitute the division by a bit shift

Algorithm optimization

- Do we need a lot of precision in the division?
- Not really, we could substitute the division by a bit shift
- It is important to:
 - Check constant values
 - Check Underflow and Overflow conditions
 - Minimize the introduced error

Algorithm optimization

- Do we need a lot of precision in the division?
- Not really, we could substitute the division by a bit shift
- It is important to:
 - Check constant values
 - Check Underflow and Overflow conditions
 - Minimize the introduced error
- By removing the division, you can gain from ~800 to ~580 cycles/pixel (1.38x)

Algorithm optimization

- With this optimization, what happened with the players:
area, performance, power consumption, time-to-market?

Algorithm optimization

- With this optimization, what happened with the players:
area, performance, power consumption, time-to-market?
- Can we optimize more?

Nios II Reference Guide

Table 60. Hardware Multiply and Divide Details for the Nios II/f Core

ALU Option	Hardware Details	Cycles per Instruction	Result Latency Cycles	Supported Instructions
No hardware multiply or divide	Multiply and divide instructions generate an exception	–	–	None
Logic elements	ALU includes 32 x 4-bit multiplier	11	+2	mul, muli
<i>continued...</i>				
32-bit multiplier	ALU includes 32 x 32-bit multiplier	1	+2	mul, muli, mulxss, mulxsu, mulxuu
16-bit multiplier	ALU includes 3 16 x 16-bit multiplier	1	+2	mul, muli
16-bit multiplier	ALU includes 4 16 x 16-bit multiplier	2	+2	mul, muli, mulxss, mulxsu, mulxuu
Hardware divide	ALU includes SRT Radix-2 divide circuit	35	+2	div, divu

Algorithm optimization

- With this optimization, what happened with the players:
area, performance, power consumption, time-to-market?
- Can we optimize more?
 - Using optimized multipliers?

Algorithm optimization

- With this optimization, what happened with the players: area, performance, power consumption, time-to-market?
- Can we optimize more?
 - Using optimized multipliers?
 - Other parts of the code?

Other optimizations?

$$\text{Grayscale} \approx 0.33 \cdot R + 0.56 \cdot G + 0.11 \cdot B$$

```
void conv_grayscale(void *picture,
                   int width,
                   int height) {
    int x,y,gray;
    unsigned short *pixels = (unsigned short *)picture , rgb;
    grayscale_width = width;
    grayscale_height = height;
    if (grayscale_array != NULL)
        free(grayscale_array);
    grayscale_array = (unsigned char *) malloc(width*height);
    for (y = 0 ; y < height ; y++) {
        for (x = 0 ; x < width ; x++) {
            rgb = pixels[y*width+x];
            gray = (((rgb>>11)&0x1F)<<3)*21; // red part
            gray += (((rgb>>5)&0x3F)<<2)*72; // green part
            gray += (((rgb>>0)&0x1F)<<3)*7; // blue part
            gray /= 100;
            IOWR_8DIRECT(grayscale_array,y*width+x,gray);
        }
    }
}
```


~800 → ~480 cycles/pixel

Algorithm optimization

- With this optimization, what happened with the players:
area, performance, power consumption, time-to-market?
- Can we optimize more?
 - Using optimized multipliers?
 - Other parts of the code?
 - Anything else?

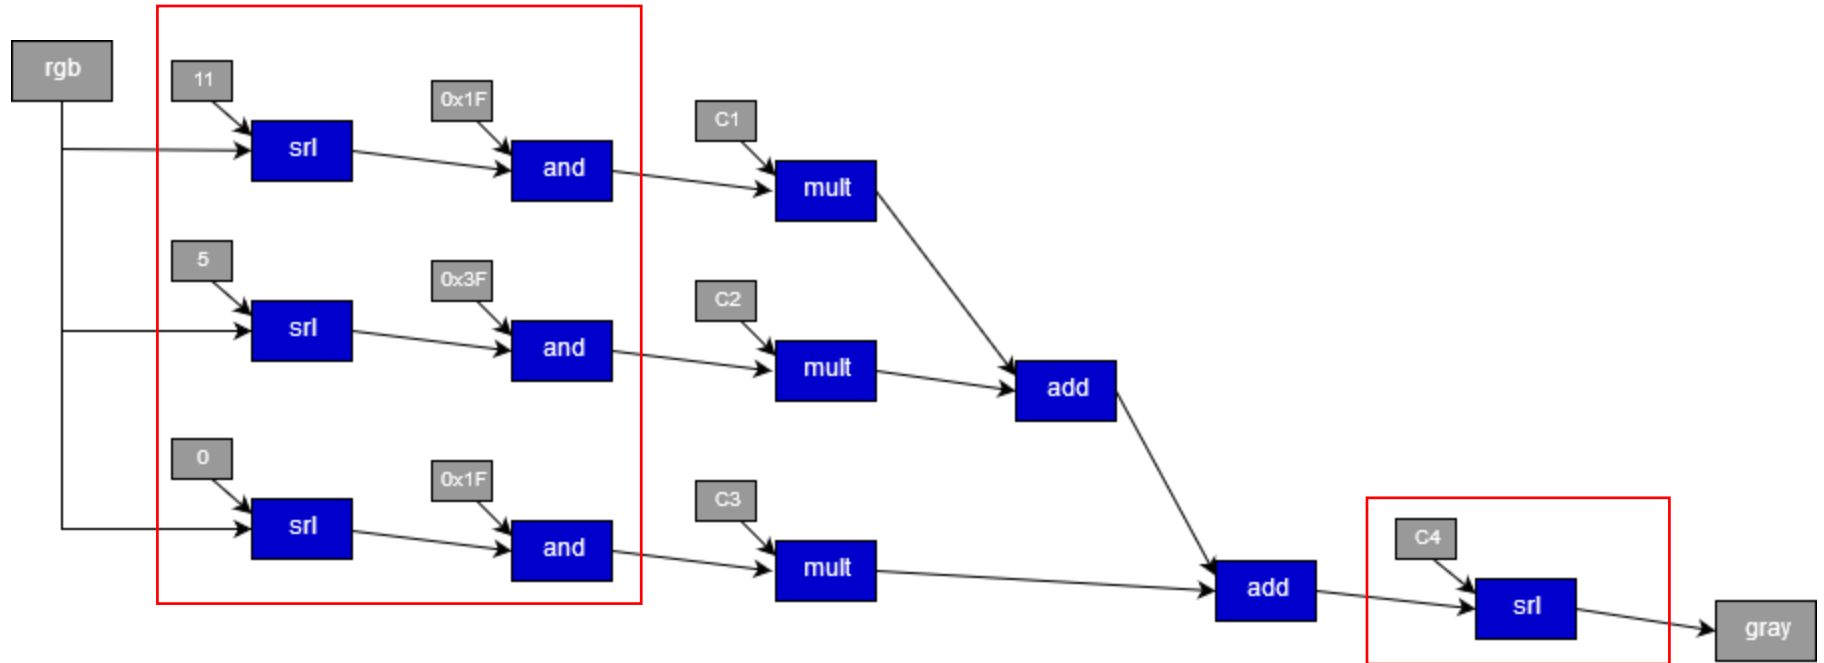
Algorithm optimization

➤ How would you implement it in HW?



Algorithm optimization

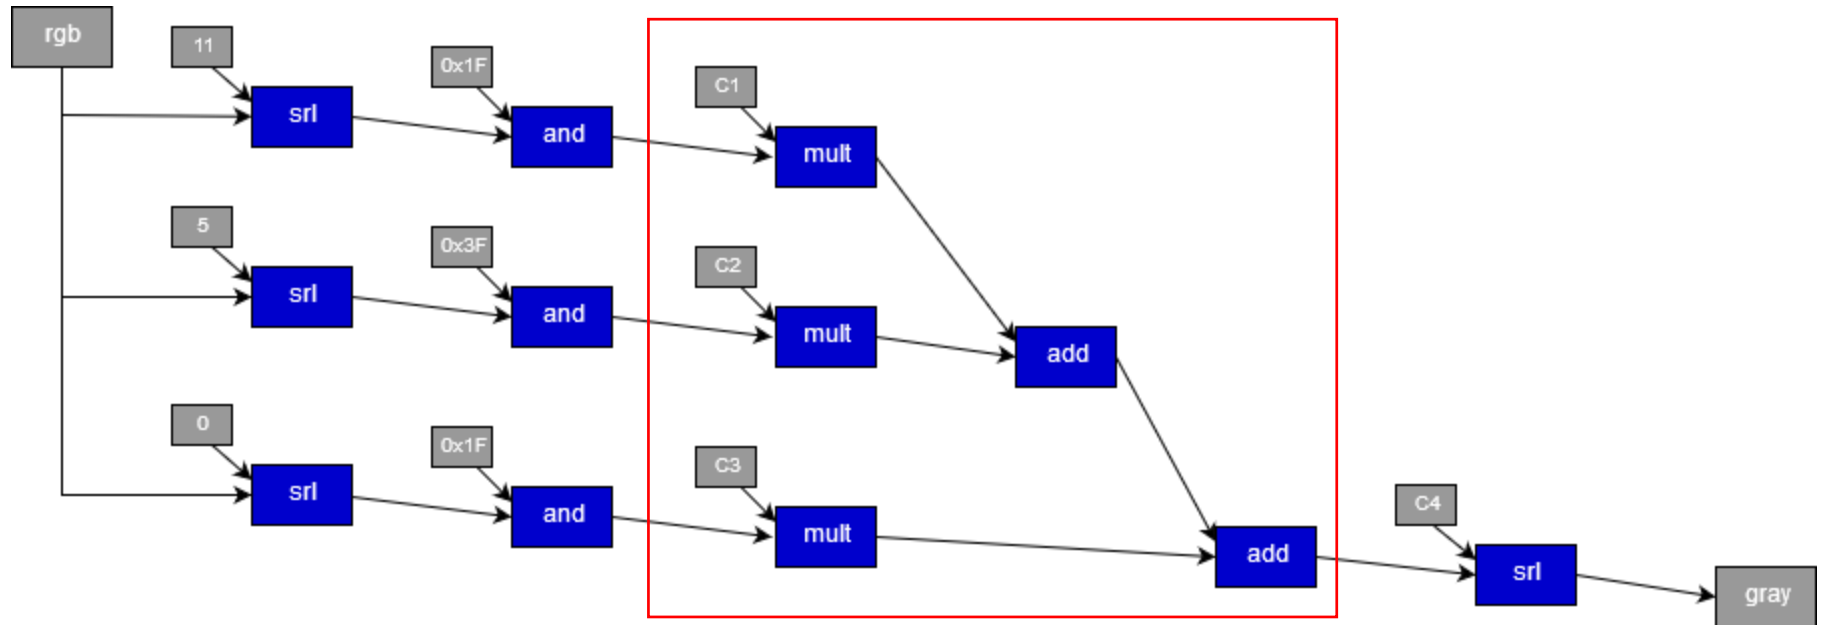
➤ How would you implement it in HW?



➤ Shifting and slicing

Algorithm optimization

➤ How would you implement it in HW?



➤ Shifting and slicing

➤ Multipliers and adders

Algorithm optimization

- How would you implement it in HW?
 - Shifting and slicing
 - Multipliers and adders
 - Parallelism

Algorithm optimization

- How would you implement it in HW?
 - Shifting and slicing
 - Multipliers and adders
 - Parallelism
- In SW:
 - 1 LOAD + 1 OPERATION + 1 STORE
 - + loop control

Algorithm optimization

- How would you implement it in HW?
 - Shifting and slicing
 - Multipliers and adders
 - Parallelism
- In SW:
 - 1 LOAD + 1 OPERATION + 1 STORE
 - + loop control
- What happened with the players: area, performance, power consumption, time-to-market?